

Tyler Freemont

AI Infrastructure Engineer — MCP Servers · RAG · Local LLM Deployment

Python · TypeScript · C++ · Self-hosted inference · EN/JA bilingual

Osaka, Japan · tfreemont0712@gmail.com · 080-2345-4653 · tylerfreemont0712.github.io · github.com/tylerfreemont0712

SUMMARY

I build production LLM infrastructure for teams that cannot send their data to external providers: self-hosted inference (Ollama, llama.cpp) behind private API endpoints, Model Context Protocol servers, RAG pipelines and multi-agent tooling in Python and TypeScript. Backed by an infrastructure career in carrier-scale network automation and AWS. Native English speaker, JLPT N1 Japanese, based in Japan.

CORE SKILLS

AI / LLM systems: Local LLM deployment (Ollama, llama.cpp) · Model Context Protocol (MCP) servers · RAG / retrieval pipelines · LangChain · Multi-agent orchestration · Tool-calling & function calling · Prompt engineering · Model evaluation (RLHF) · Benchmarking · Qwen, DeepSeek, Gemma, gpt-oss

Programming languages: Proficient: Python, TypeScript, JavaScript, C++, C#, SQL, HTML/CSS · Working knowledge: Go, Rust

Infrastructure & backend: Linux · Docker · AWS · CI/CD · REST API design · Network automation · TCP/IP · Fault detection & auto-recovery · Git

Frontend & applications: React · PyQt6 · HTML5 Canvas · SaaS product development

Languages: English (native) · Japanese (JLPT N1) — technical documentation, client-facing meetings, and EN↔JP technical translation

CERTIFICATIONS

AWS Certified Developer – Associate · AWS Certified Solutions Architect – Associate · JLPT N1 (Japanese Language Proficiency Test, highest level) · Google IT Support Professional Certificate

EXPERIENCE

Software Engineer, AI Systems

Mar 2026 – Present

MiraX · Osaka, Japan (Remote)

- Built the company's local LLM inference layer (Ollama, llama.cpp) behind internal API endpoints, moving enterprise workloads off external model providers and keeping customer data in-house.
- Designed and hardened 3 Model Context Protocol (MCP) servers now in production use by the company's internal local LLM deployment — including a company-document server with full Japanese filesystem and Unicode handling, and a bilingual EN/JA web-search service.
- Shipped RAG retrieval, tool-calling and multi-agent workflows (LangChain) into an enterprise SaaS platform spanning video conferencing, translation and records management.

AI Evaluation Specialist (Code & Math)

Jan 2024 – Present

Outlier / Alignerr · Remote (Freelance)

- Produce RLHF preference data and detailed written technical justifications used directly in frontier-model training pipelines.

Software / Infrastructure Engineer

Sep 2023 – Feb 2026

Neighbors Inc. · Osaka, Japan

- Built Python automation tooling and internal services for enterprise network clients including Rakuten Mobile and Daikin.
- Designed and implemented a fault-detection and auto-recovery service covering 10,000+ network nodes, converting operator-driven incident recovery into an automatic path.

SELECTED PROJECTS

LLM Council — Multi-Agent Autonomous Build System

Python

- Orchestrates specialised AI agents through a shared tool layer to autonomously build working HTML5 Canvas browser games end to end.
- Implements filesystem access, web search, asset retrieval, hot-reloading and structured inter-agent messaging, with AST-level code validation and runtime tool generation.